



The Beauty and Joy of Computing

Lecture #8 Concurrency



CS10 Head TA
Michael Ball

"KOOOMEY'S LAW" – EFFICIENCY 2X EVERY 18 MO

Prof Jonathan Koomey looked at 6 decades of data and found that energy efficiency of computers doubles roughly every 18 months. This is even more relevant as battery-powered devices become more popular. Restated, it says that for a fixed computing load, the amount of battery you need drops by half every 18 months. This was true before transistors!



www.technologyreview.com/computing/38548/



Concurrency: A Definition

Concurrency: A property of computer systems in which several computations are executing simultaneously, and potentially interacting with each other.



UC Berkeley "The Beauty and Joy of Computing" : Concurrency (2)



Concurrency is Everywhere!

Examples:

- Mouse cursor movement while Snap! calculates.
- Screen clock advances while typing in a text.
- Busy cursor spins while browser connects to server, waiting for response
- Walking while chewing gum



UC Berkeley "The Beauty and Joy of Computing" : Concurrency (3)



Concurrency & Parallelism

Intra-computer

- Today's lecture
- Multiple computing "helpers" are cores within one machine
- Aka "multi-core"
 - Although GPU parallism is also "intra-computer"

Inter-computer

- Future lecture
- Multiple computing "helpers" are different machines
- Aka "distributed computing"
 - Grid & cluster computing




UC Berkeley "The Beauty and Joy of Computing" : Concurrency (4)



Anatomy: 5 components of any Computer

John von Neumann
invented this
architecture



- a) Control
- b) Datapath
- c) Memory
- d) Input
- e) Output


What causes the most headaches for SW and HW designers with multi-core computing?



UC Berkeley "The Beauty and Joy of Computing" : Concurrency (5)



But what is INSIDE a Processor?

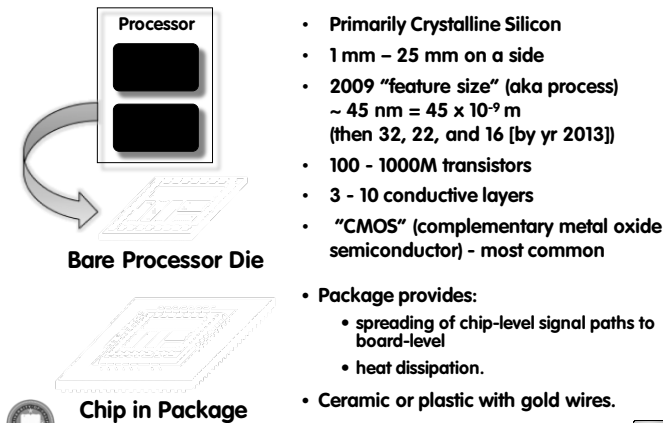


UC Berkeley "The Beauty and Joy of Computing" : Concurrency (6)





But what is INSIDE a Processor?





UC Berkeley “The Beauty and Joy of Computing”: Concurrency (7)



en.wikipedia.org/wiki/Moore's Law

Moore's Law

Predicts: 2X Transistors / chip every 2 years

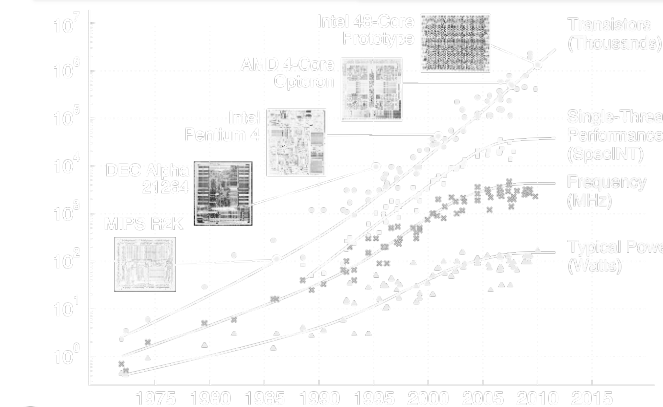


Gordon Moore
Intel Cofounder
B.S. Cal 1950!

UC Berkeley “The Beauty and Joy of Computing”: Concurrency (8)



Moore's Law and related curves

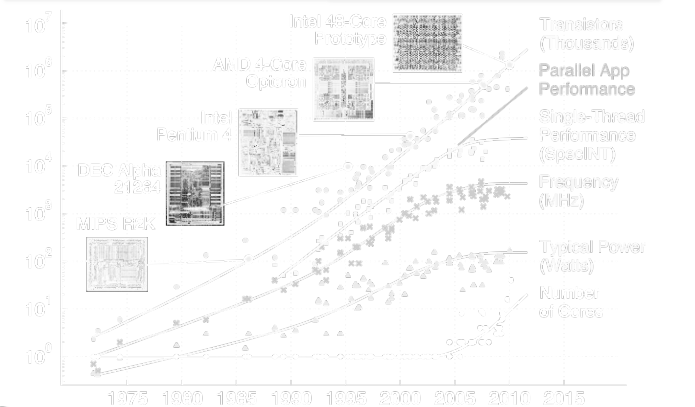


Data partially collected by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond

UC Berkeley “The Beauty and Joy of Computing”: Concurrency (9)



Moore's Law and related curves




Data partially collected by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond

UC Berkeley “The Beauty and Joy of Computing”: Concurrency (10)



Power Density Prediction circa 2000




Source: S. Borkar (Intel)

UC Berkeley “The Beauty and Joy of Computing”: Concurrency (11)



Background: Threads

- A **Thread** stands for “thread of execution”, is a single stream of instructions
 - A program / process can split, or fork itself into separate threads, which can (in theory) execute simultaneously.
 - An easy way to describe/think about parallelism
- A single CPU can execute many threads by **Time Division Multiplexing**
- Multithreading** is running multiple threads through the same hardware



UC Berkeley “The Beauty and Joy of Computing”: Concurrency (12)






en.wikipedia.org/wiki/Amdahl's_law

Speedup Issues : Amdahl's Law

- Applications can almost never be completely parallelized; some serial code remains



- s is serial fraction of program, P is # of cores (was processors)

Amdahl's law:

$$\text{Speedup}(P) = \text{Time}(1) / \text{Time}(P)$$

$$\leq 1 / (s + [(1-s) / P]), \text{ and as } P \rightarrow \infty$$

$$\leq 1 / s$$

- Even if the parallel portion of your application speeds up perfectly, your performance may be limited by the sequential portion



UC Berkeley "The Beauty and Joy of Computing" : Concurrency (13)



Speedup Issues : Overhead

Even assuming no sequential portion, there's...

- Time to think how to divide the problem up
- Time to hand out small "work units" to workers
- All workers may not work equally fast
- Some workers may fail
- There may be contention for shared resources
- Workers could overwriting each others' answers
- You may have to wait until the last worker returns to proceed (the slowest / weakest link problem)
- There's time to put the data back together in a way that looks as if it were done by one



UC Berkeley "The Beauty and Joy of Computing" : Concurrency (14)



Life in a multi-core world...

- This "sea change" to multi-core parallelism means that the computing community has to rethink:

- Languages
- Architectures
- Algorithms
- Data Structures
- All of the above



UC Berkeley "The Beauty and Joy of Computing" : Concurrency (15)



en.wikipedia.org/wiki/Concurrent_computing


But parallel programming is hard!

- What if two people were calling withdraw at the same time?

- E.g., balance=100 and two withdraw 75 each
- Can anyone see what the problem *could* be?
- This is a race condition

- In most languages, this is a problem.

- In Snap!, the system doesn't let two of these run at once.



UC Berkeley "The Beauty and Joy of Computing" : Concurrency (16)



"Non-Deterministic" Parallel Code

- Two (or more) scripts are running at the same time, BUT we don't know what order they will be run in!
- Each individual script runs its blocks in order, but the processor (Snap!) will swap between running script A and script B.





UC Berkeley "The Beauty and Joy of Computing" : Concurrency (17)



How Many Possible Outputs?

- We want this code to draw a cute winky-face, but there's a problem with parallelizing it!
- How many possible outputs can we have?

- 1
- 3
- 4
- 7
- 8



UC Berkeley "The Beauty and Joy of Computing" : Concurrency (18)





Another concurrency problem ... deadlock!

- Two people need to draw a graph but there is only one pencil and one ruler.
 - One grabs the pencil
 - One grabs the ruler
 - Neither release what they hold, waiting for the other to release
- Livelock also possible
 - Movement, no progress



Summary

- "Sea change" of computing because of inability to cool CPUs means we're now in multi-core world
- This brave new world offers lots of potential for innovation by computing professionals, but challenges persist

